

Ryan Long
DSC 680-T301
Project 2

```
In [1]: #import Libraries
import pandas as pd
import numpy as np
from pandas_profiling import ProfileReport
import seaborn as sns
```

Data Import & Clean

```
In [2]: # https://www.kaggle.com/datasets/dhanushnarayananr/credit-card-fraud
data = pd.read_csv('card_transdata.csv')
```

EDA

```
In [3]: profile = ProfileReport(data,title="Pandas Profiling Report",explorative=True)
```

In [4]: profile

Summarize dataset:	30/30 [00:38<00:00, 2.15it/s,
100%	Completed]
Generate report structure:	1/1 [00:02<00:00,
100%	2.79s/it]
Render HTML: 100%	1/1 [00:00<00:00, 1.26it/s]

Overview

Dataset statistics

Number of variables	8
Number of observations	1000000
Missing cells	0
Missing cells (%)	0.0%
Duplicate rows	0
Duplicate rows (%)	0.0%
Total size in memory	61.0 MiB
Average record size in memory	64.0 B

Variable types

Numeric	3
Categorical	5

Alerts

distance_from_home is highly correlated with repeat_retailer	High correlation
repeat_retailer is highly correlated with distance_from_home	High correlation
distance_from_home is highly skewed ($\gamma_1 = 20.2397331$)	Skewed

Out[4]:

In [5]: `profile.to_file("Project_2_EDA.html")`

Export report to file:

100%

1/1 [00:00<00:00,

37.03it/s]

In [6]: *#dataset appears to be pre-processed*
data

Out[6]:

	distance_from_home	distance_from_last_transaction	ratio_to_median_purchase_price	repeat
0	57.877857	0.311140	1.945940	
1	10.829943	0.175592	1.294219	
2	5.091079	0.805153	0.427715	
3	2.247564	5.600044	0.362663	
4	44.190936	0.566486	2.222767	
...
999995	2.207101	0.112651	1.626798	
999996	19.872726	2.683904	2.778303	
999997	2.914857	1.472687	0.218075	
999998	4.258729	0.242023	0.475822	
999999	58.108125	0.318110	0.386920	

1000000 rows × 8 columns



In [7]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000000 entries, 0 to 999999
Data columns (total 8 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   distance_from_home                    1000000 non-null float64
1   distance_from_last_transaction        1000000 non-null float64
2   ratio_to_median_purchase_price        1000000 non-null float64
3   repeat_retailer                       1000000 non-null float64
4   used_chip                             1000000 non-null float64
5   used_pin_number                       1000000 non-null float64
6   online_order                          1000000 non-null float64
7   fraud                                 1000000 non-null float64
dtypes: float64(8)
memory usage: 61.0 MB
```

In [8]: `data.describe()`

Out[8]:

	distance_from_home	distance_from_last_transaction	ratio_to_median_purchase_price	repeat
count	1000000.000000	1000000.000000	1000000.000000	1000000.000000
mean	26.628792	5.036519	1.824182	1.824182
std	65.390784	25.843093	2.799589	2.799589
min	0.004874	0.000118	0.004399	0.004399
25%	3.878008	0.296671	0.475673	0.475673
50%	9.967760	0.998650	0.997717	0.997717
75%	25.743985	3.355748	2.096370	2.096370
max	10632.723672	11851.104565	267.802942	267.802942

Modeling

In [9]: `# model library/package imports`
`from xgboost import XGBClassifier`
`from sklearn.model_selection import train_test_split`
`from sklearn.metrics import accuracy_score, confusion_matrix, classification_report`

In [10]: `# isolate target and features`
`X = data.iloc[:, 0:7].columns`
`Y = data.iloc[:, 7:].columns`

In [11]: `X`

Out[11]: Index(['distance_from_home', 'distance_from_last_transaction',
'ratio_to_median_purchase_price', 'repeat_retailer', 'used_chip',
'used_pin_number', 'online_order'],
dtype='object')

In [12]: `Y`

Out[12]: Index(['fraud'], dtype='object')

In [13]: `X = data[X]`
`Y = data[Y]`

```
In [14]: # split data into train, test, and validation sets
train_ratio = 0.65
test_ratio = 0.20
validation_ratio = 0.15

X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=test_ratio)
X_train, X_valid, y_train, y_valid = train_test_split(X_train, y_train, test_size=
```

```
In [15]: print(X_train.shape)
print(X_test.shape)
print(X_valid.shape)
```

```
(658823, 7)
(200000, 7)
(141177, 7)
```

```
In [16]: print("Train fraud, not fraud:",(y_train['fraud'] != 0).sum(),(y_train['fraud'] == 0).sum())
print("Test fraud, not fraud:",(y_test['fraud'] != 0).sum(),(y_test['fraud'] == 0).sum())
print("Validation fraud, not fraud:",(y_valid['fraud'] != 0).sum(),(y_valid['fraud'] == 0).sum())
```

```
Train fraud, not fraud: 57509 601314
Test fraud, not fraud: 17541 182459
Validation fraud, not fraud: 12353 128824
```

```
In [17]: # fit model to training data
model = XGBClassifier()
model.fit(X_train, y_train)
```

```
Out[17]: XGBClassifier
XGBClassifier(base_score=0.5, booster='gbtree', callbacks=None,
              colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1,
              early_stopping_rounds=None, enable_categorical=False,
              eval_metric=None, gamma=0, gpu_id=-1, grow_policy='depthwise',
              importance_type=None, interaction_constraints='',
              learning_rate=0.300000012, max_bin=256, max_cat_to_onehot=4,
              max_delta_step=0, max_depth=6, max_leaves=0, min_child_weight=
1,
              missing=nan, monotone_constraints='()', n_estimators=100,
              n_jobs=0, num_parallel_tree=1, predictor='auto', random_state=
```

Model Evaluation

```
In [18]: # make predictions for test data
y_pred = model.predict(X_test)
predictions = [value for value in y_pred]
```


```
In [19]: # evaluate predictions
accuracy = accuracy_score(y_test, predictions)
print(accuracy)
```

0.999985

```
In [20]: # confusion matrix for test set
print(confusion_matrix(y_test, predictions))
cfm = confusion_matrix(y_test, predictions)
sns.heatmap(cfm,annot=True,fmt="d")
```

```
[[182458    1]
 [     2 17539]]
```

Out[20]: <AxesSubplot:>



```
In [21]: # classification for test set
print(classification_report(y_test, predictions))
```


	precision	recall	f1-score	support
0.0	1.00	1.00	1.00	182459
1.0	1.00	1.00	1.00	17541
accuracy			1.00	200000
macro avg	1.00	1.00	1.00	200000
weighted avg	1.00	1.00	1.00	200000

```
In [22]: # make predictions for validation data
y_pred = model.predict(X_valid)
predictions = [value for value in y_pred]
```

```
In [23]: # confusion matrix and classification for validation set
print(confusion_matrix(y_valid, predictions))
cfm = confusion_matrix(y_valid, predictions)
sns.heatmap(cfm, annot=True, fmt="d")
```

```
[[128824    0]
 [    0 12353]]
```

Out[23]: <AxesSubplot:>



```
In [24]: # classification for validation set
print(classification_report(y_valid, predictions))
```

	precision	recall	f1-score	support
0.0	1.00	1.00	1.00	128824
1.0	1.00	1.00	1.00	12353
accuracy			1.00	141177
macro avg	1.00	1.00	1.00	141177
weighted avg	1.00	1.00	1.00	141177

```
In [25]: print("Accuracy on training set: {:.3f}".format(model.score(X_train, y_train)))
print("Accuracy on validation set: {:.3f}".format(model.score(X_valid, y_valid)))
```

```
Accuracy on training set: 1.000
Accuracy on validation set: 1.000
```